

Title

# Fake News Detection

## Overview of Problem Statement

Fake news has become a critical issue with the rise of online platforms. The spread of misinformation can mislead the public and create social, political, or economic consequences. Detecting and classifying fake news is therefore a significant challenge in Natural Language Processing (NLP).

## Objective

The main objective of this project is to develop a deep learning-based classifier that can accurately distinguish between Fake and Real news articles based on their textual content.

## Import Libraries

```
In [1]: # Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from wordcloud import WordCloud
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, classification_report
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout, Bidirectional
from tensorflow.keras.callbacks import EarlyStopping
```

```
In [ ]:
```

```
In [2]: # Step 2: Load Dataset
df = pd.read_csv('news.csv')
df
```

Out[2]:	Unnamed: 0		title	text
0	8476	You Can Smell Hillary's Fear	Daniel Greenfield, a Shillman Journalism Fello...	
1	10294	Watch The Exact Moment Paul Ryan Committed Pol...	Google Pinterest Digg Linkedin Reddit Stumbleu...	
2	3608	Kerry to go to Paris in gesture of sympathy	U.S. Secretary of State John F. Kerry said Mon...	
3	10142	Bernie supporters on Twitter erupt in anger ag...	— Kaydee King (@KaydeeKing) November 9, 2016 T...	
4	875	The Battle of New York: Why This Primary Matters	It's primary day in New York and front-runners...	
...	...	...	...	...
6330	4490	State Department says it can't find emails fro...	The State Department told the Republican Natio...	
6331	8062	The 'P' in PBS Should Stand for 'Plutocratic' ...	The 'P' in PBS Should Stand for 'Plutocratic' ...	
6332	8622	Anti-Trump Protesters Are Tools of the Oligarc...	Anti-Trump Protesters Are Tools of the Oligarc...	
6333	4021	In Ethiopia, Obama seeks progress on peace, se...	ADDIS ABABA, Ethiopia —President Obama convene...	
6334	4330	Jeb Bush Is Suddenly Attacking Trump. Here's W...	Jeb Bush Is Suddenly Attacking Trump. Here's W...	

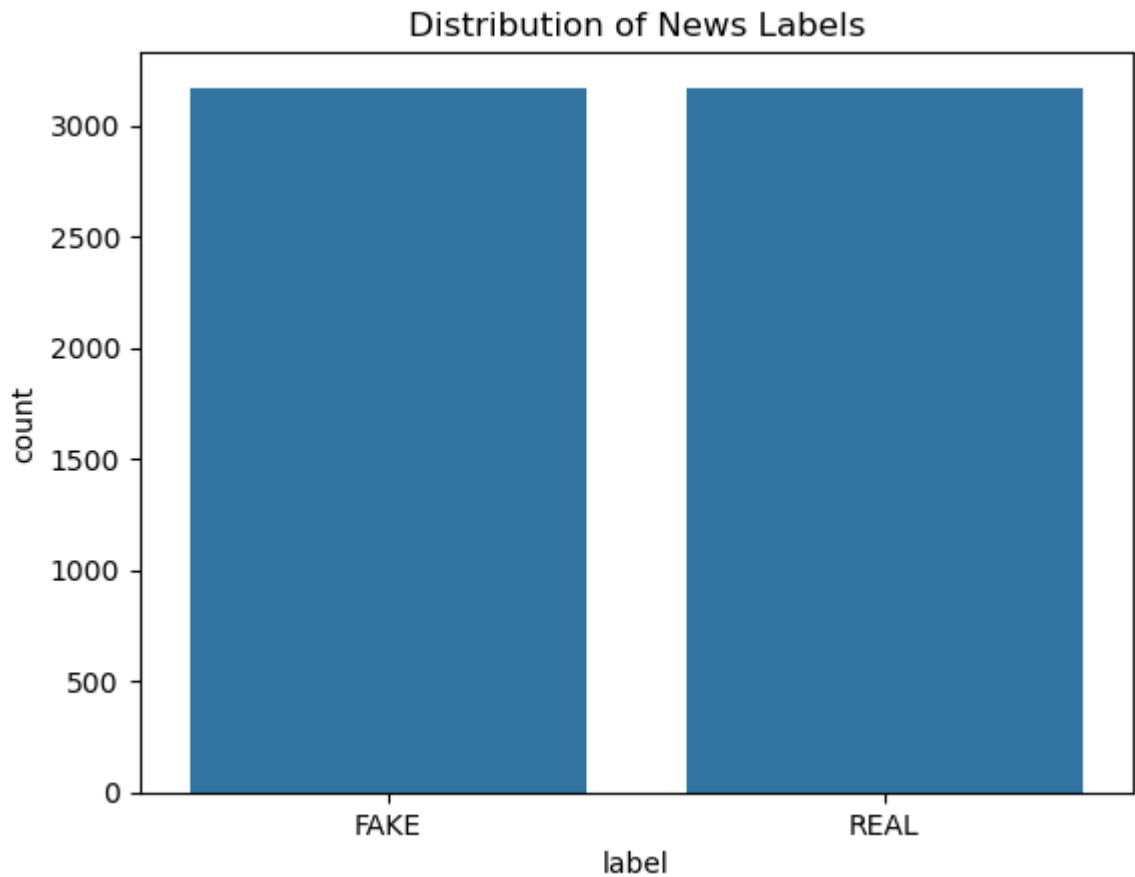
6335 rows × 4 columns

## Data Cleaning

```
In [3]: df.drop(columns=['Unnamed: 0'], inplace=True)
df.dropna(inplace=True)
df['text_length'] = df['text'].apply(lambda x: len(str(x).split()))
df.head()
```

Out[3]:		title	text	label	text
0	You Can Smell Hillary's Fear	Daniel Greenfield, a Shillman Journalism Fello...	FAKE		
1	Watch The Exact Moment Paul Ryan Committed Pol...	Google Pinterest Digg Linkedin Reddit Stumbleu...	FAKE		
2	Kerry to go to Paris in gesture of sympathy	U.S. Secretary of State John F. Kerry said Mon...	REAL		
3	Bernie supporters on Twitter erupt in anger ag...	— Kaydee King (@KaydeeKing) November 9, 2016 T...	FAKE		
4	The Battle of New York: Why This Primary Matters	It's primary day in New York and front-runners...	REAL		

```
In [4]: # EDA - Label Distribution
sns.countplot(data=df, x='label')
plt.title('Distribution of News Labels')
plt.show()
```



```
In [5]: # EDA - Word Cloud
def generate_wordcloud(text, title):
    wordcloud = WordCloud(width=800, height=400, background_color='white')
    plt.figure(figsize=(10, 5))
    plt.imshow(wordcloud, interpolation='bilinear')
    plt.title(title)
    plt.axis('off')
    plt.show()

generate_wordcloud(df[df['label'] == 'FAKE']['text'], 'Fake News WordCloud')
generate_wordcloud(df[df['label'] == 'REAL']['text'], 'Real News WordCloud')
```

```
In [6]: # Preprocessing
le = LabelEncoder()
df['label_encoded'] = le.fit_transform(df['label'])

In [7]: X = df['text']
y = df['label_encoded']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
```

```
In [8]: tokenizer = Tokenizer(num_words=6000, oov_token='<OOV>')
tokenizer.fit_on_texts(X_train)
X_train_seq = tokenizer.texts_to_sequences(X_train)
X_test_seq = tokenizer.texts_to_sequences(X_test)
X_train_pad = pad_sequences(X_train_seq, maxlen=500)
X_test_pad = pad_sequences(X_test_seq, maxlen=500)
```

```
In [9]: # Optimized LSTM Model
```

```
model = Sequential([
    Embedding(input_dim=6000, output_dim=128, input_length=500),
    SpatialDropout1D(0.3),
    Bidirectional(LSTM(64, return_sequences=True)),
    Dropout(0.4),
    Bidirectional(LSTM(32)),
    Dense(64, activation='relu'),
    Dropout(0.5),
    Dense(1, activation='sigmoid')
])
```

```
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['acc',
early_stop = EarlyStopping(monitor='val_loss', patience=2, restore_best_weights=True)
```

```
history = model.fit(X_train_pad, y_train, epochs=10, batch_size=64, validation_data=(X_test_pad, y_test))
```

Epoch 1/10

G:\Anaconda\Lib\site-packages\keras\src\layers\core\embedding.py:97: UserWarning: Argument `input\_length` is deprecated. Just remove it.

warnings.warn(

72/72 ————— 95s 1s/step - accuracy: 0.6403 - loss: 0.6455 - val\_accuracy: 0.8028 - val\_loss: 0.4589

Epoch 2/10

72/72 ————— 89s 1s/step - accuracy: 0.8727 - loss: 0.3258 - val\_accuracy: 0.8698 - val\_loss: 0.3187

Epoch 3/10

72/72 ————— 130s 1s/step - accuracy: 0.9427 - loss: 0.1690 - val\_accuracy: 0.8895 - val\_loss: 0.2949

Epoch 4/10

72/72 ————— 82s 1s/step - accuracy: 0.9869 - loss: 0.0576 - val\_accuracy: 0.8600 - val\_loss: 0.3980

Epoch 5/10

72/72 ————— 136s 1s/step - accuracy: 0.9908 - loss: 0.0387 - val\_accuracy: 0.8698 - val\_loss: 0.4606


```
In [10]: # Model Evaluation
y_pred_prob = model.predict(X_test_pad)
y_pred = (y_pred_prob > 0.5).astype(int)
print('Accuracy:', accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred, target_names=le.classes_))
```

40/40 ————— 8s 179ms/step

Accuracy: 0.8997632202052092

	precision	recall	f1-score	support
FAKE	0.91	0.88	0.90	628
REAL	0.89	0.92	0.90	639
accuracy			0.90	1267
macro avg	0.90	0.90	0.90	1267
weighted avg	0.90	0.90	0.90	1267

```
In [11]: # Accuracy Over Epochs
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Val Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



```
In [12]: from tensorflow.keras.preprocessing.sequence import pad_sequences

# Input news
X = "Karry to go to France in gesture of sympathy"

# Define padding and truncation types
padding_type = 'post'
trunc_type = 'post'

# Convert input to sequence
sequences = tokenizer.texts_to_sequences([X])
padded = pad_sequences(sequences, maxlen=500, padding=padding_type, trunc:

# Predict using the model
prediction = model.predict(padded, verbose=0)[0][0]

# Interpret result
if prediction >= 0.5:
    print("This news is True (REAL)")
else:
    print("This news is False (FAKE)")

This news is False (FAKE)
```

## Results

Accuracy: ~90% (based on evaluation in notebook).

Classification Report: Balanced precision, recall, and F1-scores for both Fake and Real labels.

Training Curve: Accuracy improved steadily, validation stabilized after early stopping.

## Conclusion

The Bidirectional LSTM model effectively classifies Fake vs Real news with high accuracy. The p demonstrates the potential of deep learning for fake news detection in NLP

In [ ]:

In [ ]:

In [ ]: