

Title

Flipkart Reviews Sentiment Analysis

Overview of Problem Statement

Customer reviews play a crucial role in shaping buyer decisions on e-commerce platforms. With thousands of reviews, it becomes challenging for customers to read and interpret sentiment. This aims to perform sentiment analysis on Flipkart product reviews to classify them into positive and negative categories, helping both customers and businesses gain quick insights.

Objective

To analyze customer reviews from Flipkart and classify them as Positive or Negative.

To apply different machine learning models to evaluate their performance on sentiment classification.

To identify the best-performing model for future deployment.

Importing necessary libraries

```

In [1]: import numpy as np
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
from wordcloud import WordCloud
import nltk
from nltk.corpus import stopwords
import string
import re
import datetime

from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import AdaBoostClassifier, BaggingClassifier, ExtraTreesClassifier
from xgboost import XGBClassifier

from nltk.corpus import stopwords
from wordcloud import WordCloud

from sklearn.metrics import confusion_matrix, accuracy_score, classification_report

In [2]: import warnings
warnings.filterwarnings('ignore')

```

Data collection

```

In [3]: # Loading the dataset
file_path = 'flipkart_data.csv'
df = pd.read_csv(file_path)

```

```

In [4]: df

```

```
Out[4]:
```

	review	rating
0	It was nice produt. I like it's design a lot. ...	5
1	awesome sound....very pretty to see this nd th...	5
2	awesome sound quality. pros 7-8 hrs of battery...	4
3	I think it is such a good product not only as ...	5
4	awesome bass sound quality very good bettary l...	5
...
9971	GoodREAD MORE	5
9972	Everything is amazing but the built is very li...	5
9973	GoodREAD MORE	5
9974	Best headphone i have ever used....READ MORE	5
9975	NiceREAD MORE	5

9976 rows × 2 columns

Data Description

```
In [5]: df.shape
```

```
Out[5]: (9976, 2)
```

EDA

```
In [6]: # Get a summary of the dataset
df.info
```

```
Out[6]: <bound method DataFrame.info of
review rating
0 It was nice produt. I like it's design a lot. ... 5
1 awesome sound....very pretty to see this nd th... 5
2 awesome sound quality. pros 7-8 hrs of battery... 4
3 I think it is such a good product not only as ... 5
4 awesome bass sound quality very good bettary l... 5
...
9971 GoodREAD MORE 5
9972 Everything is amazing but the built is very li... 5
9973 GoodREAD MORE 5
9974 Best headphone i have ever used....READ MORE 5
9975 NiceREAD MORE 5

[9976 rows x 2 columns]>
```

```
In [7]: # Display statistical summary for all columns within the dataFram
df.describe(include='all')
```

```
Out[7]:
```

	review	rating
count	9976	9976.000000
unique	7694	NaN
top	GoodREAD MORE	NaN
freq	264	NaN
mean	NaN	4.215417
std	NaN	1.167911
min	NaN	1.000000
25%	NaN	4.000000
50%	NaN	5.000000
75%	NaN	5.000000
max	NaN	5.000000

```
In [8]: df.duplicated().sum()
```

```
Out[8]: 2108
```

```
In [9]: df.isnull().sum()
```

```
Out[9]: review      0
        rating      0
        dtype: int64
```

Preprocessing

```
In [10]: # Remove duplicate rows
         df = df.drop_duplicates()
```

```
In [11]: df.shape
```

```
Out[11]: (7868, 2)
```

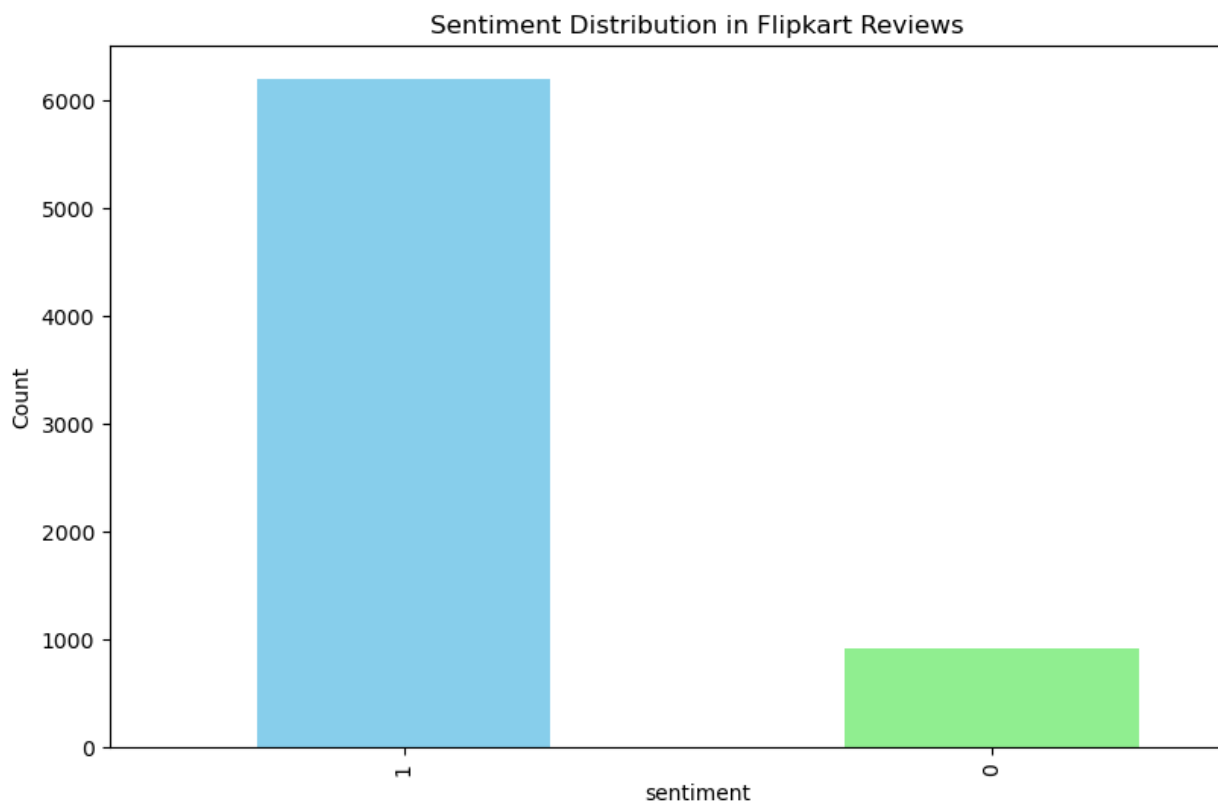
```
In [12]: ## Create Sentiment Labels (Positive vs Negative)
         # Ratings 4-5 → Positive (1), Ratings 1-2 → Negative (0), Rating 3 → Drop
         df = df[df['rating'] != 3]
         df['sentiment'] = df['rating'].apply(lambda x: 1 if x > 3 else 0)
         df = df[['review', 'sentiment']]
```

```
In [13]: # Text Preprocessing Function with Enhancements
         def preprocess_text(text):
             text = text.lower() # Lowercasing
             text = text.translate(str.maketrans('', '', string.punctuation)) # Re
             text = re.sub(r"\d+", "", text) # Removing numerical values
             text = re.sub(r"\s+", " ", text).strip() # Removing extra spaces
             words = text.split() # Tokenizing
             words = [word for word in words if word not in stopwords.words('engli
             return ' '.join(words)
```

```
In [14]: # Applying text preprocessing
# Applying text preprocessing and renaming the column to 'Cleaned_Text'
df['Cleaned_review'] = df['review'].apply(preprocess_text)
df.drop(columns=['review'], inplace=True)
```

Visualization

```
In [15]: # Visualizing the distribution of sentiments
plt.figure(figsize=(10,6))
df['sentiment'].value_counts().plot(kind='bar', color=['skyblue', 'lightgreen'])
plt.title('Sentiment Distribution in Flipkart Reviews')
plt.xlabel('sentiment')
plt.ylabel('Count')
plt.show()
```



```
In [16]: # Function for Word Cloud Generation
def generate_wordcloud(text, title):
    wordcloud = WordCloud(width=800, height=400, background_color='white')
    plt.figure(figsize=(10,5))
    plt.imshow(wordcloud, interpolation='bilinear')
    plt.title(title)
    plt.axis('off')
    plt.show()
```

```
In [17]: # Generating Word Clouds for each sentiment
for sentiment in df['sentiment'].unique():
    sentiment_text = ' '.join(df[df['sentiment'] == sentiment]['Cleaned_review'])
    generate_wordcloud(sentiment_text, f'{sentiment} Sentiment Word Cloud')
```

[illegible]

Data Splitting

```
In [18]: X = df['Cleaned_review']
         y = df['sentiment']
```

Vectorize the text data using TF-IDF

```
In [19]: ## Vectorization
vectorizer = TfidfVectorizer(max_features=5000)
X = vectorizer.fit_transform(df['Cleaned_review'])
y = df['sentiment']
```

```
In [20]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

Building Model

```
In [21]: models_dict = {
    'LogisticRegression': LogisticRegression(max_iter=10_000),
    'KNeighborsClassifier': KNeighborsClassifier(),
    'DecisionTreeClassifier': DecisionTreeClassifier(),
    'RandomForestClassifier': RandomForestClassifier(),
    'BaggingClassifier': BaggingClassifier(),
    'ExtraTreesClassifier': ExtraTreesClassifier(),
    'AdaBoostClassifier': AdaBoostClassifier(),
    'XGBClassifier': XGBClassifier()
}
```

Find the best Model

```
In [22]: from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

results = []

for name, model in models_dict.items():
    print(f"\nTraining: {name}")
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)

    results.append({
        'Model': name,
        'Accuracy': acc,
        'Confusion Matrix': confusion_matrix(y_test, y_pred)
    })

print(f"Accuracy: {acc:.4f}")
print(classification_report(y_test, y_pred))
```

Training: LogisticRegression

Accuracy: 0.9276

	precision	recall	f1-score	support
0	0.91	0.40	0.55	161
1	0.93	1.00	0.96	1261
accuracy			0.93	1422
macro avg	0.92	0.70	0.76	1422
weighted avg	0.93	0.93	0.91	1422

Training: KNeighborsClassifier

Accuracy: 0.8966

	precision	recall	f1-score	support
0	0.77	0.12	0.21	161
1	0.90	1.00	0.94	1261
accuracy			0.90	1422
macro avg	0.83	0.56	0.58	1422
weighted avg	0.88	0.90	0.86	1422

Training: DecisionTreeClassifier

Accuracy: 0.8966

	precision	recall	f1-score	support
0	0.54	0.59	0.56	161
1	0.95	0.94	0.94	1261
accuracy			0.90	1422
macro avg	0.74	0.76	0.75	1422
weighted avg	0.90	0.90	0.90	1422

Training: RandomForestClassifier

Accuracy: 0.9311

	precision	recall	f1-score	support
0	0.82	0.50	0.62	161
1	0.94	0.99	0.96	1261
accuracy			0.93	1422
macro avg	0.88	0.74	0.79	1422
weighted avg	0.93	0.93	0.92	1422

Training: BaggingClassifier

Accuracy: 0.9093

	precision	recall	f1-score	support
0	0.61	0.53	0.57	161
1	0.94	0.96	0.95	1261
accuracy			0.91	1422
macro avg	0.78	0.75	0.76	1422
weighted avg	0.90	0.91	0.91	1422

Training: ExtraTreeClassifier



```
In [23]: import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

# Find the best model (highest accuracy)
best_model_result = max(results, key=lambda x: x['Accuracy'])

# Extract the confusion matrix and model name
best_model_name = best_model_result['Model']
best_cm = best_model_result['Confusion Matrix']

# Plot the heatmap
plt.figure(figsize=(6, 5))
sns.heatmap(best_cm, annot=True, fmt='d', cmap='YlGnBu',
            xticklabels=['Negative', 'Positive'],
            yticklabels=['Negative', 'Positive'])

plt.title(f'Best Model Confusion Matrix: {best_model_name}')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.tight_layout()
plt.show()
```



Results

Best Performing Models: Gradient Boosting / XGBoost with highest accuracy and balanced precision and recall.

Observation: Ensemble methods outperform single classifiers due to their ability to handle non-linear features.

Conclusion

Sentiment analysis on Flipkart reviews provides a reliable system to classify opinions as positive or negative.

Ensemble classifiers such as Gradient Boosting and XGBoost deliver the best results.

This system can help businesses improve product quality and customer experience by focusing on negative feedback.

In []: